

# ECE408 Report

## APLUSPLUS

### Milestone 1 : Get Started

#### M1.1: Run the Baseline Forward Pass

Run CPU code in rai

**Our Results as below(add elapsed time for whole m1.1.py program):**

\* Running time python m1.1.py

```
New Inference
Loading fashion-mnist data... done
Loading model... done
EvalMetric: {'accuracy': 0.8673}
11.40user 11.48system
0:21.23elapsed 107%CPU (0avgtext+0a
vgdata 1628360maxresident)k
0inputs+2624outp
uts (0major+27159minor)pagefau
lts 0swaps
```

#### M1.2/1.3: Run the Baseline GPU implementation and generate a NVPROF profile

Modified rai\_build.yml according to introduction.

Mxnet GPU layer performance results is as below:

\* Running nvprof python m1.2.py

```
New Inference
Loading fashion-mnist data... done
==310== NVPROF is profiling process 310, command: python m1.2.py
Loading model...[02:36:12] src/operator/././cudnn_algoreg-inl.h:112: Running performance tests to find the best convolution
algorithm, this can take a while... (setting env variable MXNET_CUDNN_AUTOTUNE_DEFAULT to 0 to disable)
done
EvalMetric: {'accuracy': 0.8673}
==310== Profiling application: python m1.2.py
==310== Profiling result:
Time(%)   Time    Calls   Avg     Min    Max  Name
37.01%   50.022ms      1 50.022ms 50.022ms 50.022m s void cudnn::detail::implicit_convolve_sgemm<float, int=1024,
int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>(int, int, int, float const *, int,
cudnn::detail::implicit_convolve_sgemm<float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>*,
float const *, kernel_conv_params, int, float, float, int, float const *, float const *, int, int)
28.67%   38.751ms      1 38.751ms 38.751ms 38.751ms sgemm_sm35_ldg_tn_128x8x256x16x32
```

```

14.34% 19.385ms      2 9.6923ms 459.06us 18.926ms void cudnn::detail::activation_fw_4d_kernel<float, float, int=128,
int=1, int=4, cudnn::detail::tanh_func<float>>(cudnnTensorStruct, float const *, cudnn::detail::activation_fw_4d_kernel<float,
float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>, cudnnTensorStruct*, float, cudnnTensorStruct*, int,
cudnnTensorStruct*)
10.69% 14.445ms      1 14.445ms 14.445ms 14.445ms void cudnn::detail::pooling_fw_4d_kernel<float, float,
cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>(cudnnTensorStruct, float const *,
cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>,
cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)
4.53% 6.1185ms      13 470.66us 1.5360us 4.1914ms [CUDA memcpy HtoD]
2.75% 3.7160ms      1 3.7160ms 3.7160ms 3.7160ms sgemm_sm35_ldg_tn_64x16x128x8x32
0.82% 1.1115ms      1 1.1115ms 1.1115ms 1.1115ms void mshadow::cuda::SoftmaxKernel<int=8, float,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>>(mshadow::gpu, int=2, unsigned int)
0.55% 748.43us      12 62.369us 2.0800us 378.04us void mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>,
int=2)
0.32% 433.43us      2 216.72us 16.639us 416.79us void mshadow::cuda::MapPlanKernel<mshadow::sv::plusto, int=8,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::expr::Broadcast1DExp<mshadow::Tensor<mshadow::gpu, int=1, float>, float, int=2, int=1>,
float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.29% 390.68us      1 390.68us 390.68us 390.68us sgemm_sm35_ldg_tn_32x16x64x8x16
0.02% 22.399us      1 22.399us 22.399us 22.399us void mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::expr::ReduceWithAxisExp<mshadow::red::maximum, mshadow::Tensor<mshadow::gpu,
int=3, float>, float, int=3, bool=1, int=2>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.01% 10.016us      1 10.016us 10.016us 10.016us [CUDA memcpy DtoH]
==310== API calls:
Time(%)  Time    Calls    Avg     Min     Max Name
46.76% 1.85341s   18 102.97ms 17.717us 926.36ms cudaStreamCreateWithFlags
28.75% 1.13936s   10 113.94ms 696ns 322.52ms cudaFree
20.66% 818.70ms   24 34.112ms 230.81us 811.60ms cudaMemGetInfo
3.24% 128.32ms   25 5.1329ms 5.5140us 83.313ms cudaStreamSynchronize
0.31% 12.278ms   8 1.5348ms 8.4420us 6.1684ms cudaMemcpy2DAsync
0.17% 6.5406ms   42 155.73us 10.600us 1.1805ms cudaMalloc
0.03% 1.3602ms   4 340.06us 338.87us 342.68us cuDeviceTotalMem
0.02% 862.41us  114 7.5640us 625ns 307.12us cudaEventCreateWithFlags
0.02% 848.24us  352 2.4090us 248ns 63.433us cuDeviceGetAttribute
0.01% 527.45us  23 22.932us 11.237us 89.276us cudaLaunch
0.01% 419.09us   6 69.847us 59.623us 81.884us cudaMemcpy
0.01% 233.20us   4 58.298us 35.204us 83.353us cudaStreamCreate
0.00% 102.25us   4 25.561us 17.241us 32.322us cuDeviceGetName
0.00% 78.556us  32 2.4540us 695ns 8.6440us cudaSetDevice
0.00% 70.258us  110 638ns 418ns 2.3920us cudaDeviceGetAttribute
0.00% 57.402us  147 390ns 253ns 1.1690us cudaSetupArgument
0.00% 40.865us   2 20.432us 18.308us 22.557us cudaStreamCreateWithPriority
0.00% 26.085us  23 1.1340us 492ns 3.3300us cudaConfigureCall
0.00% 24.680us  10 2.4680us 1.3040us 5.9290us cudaGetDevice
0.00% 9.0930us   1 9.0930us 9.0930us 9.0930us cudaBindTexture
0.00% 8.7240us  16 545ns 407ns 891ns cudaPeekAtLastError
0.00% 4.3310us   6 721ns 240ns 1.4230us cuDeviceGetCount
0.00% 4.1760us   2 2.0880us 1.5150us 2.6610us cudaStreamWaitEvent
0.00% 3.7180us   1 3.7180us 3.7180us 3.7180us cudaStreamGetPriority
0.00% 3.4330us   6 572ns 415ns 874ns cuDeviceGet

```

0.00%	3.4270us	2	1.7130us	1.3330us	2.0940us	cudaEventRecord
0.00%	3.2360us	2	1.6180us	1.4250us	1.8110us	cudaDeviceGetStreamPriorityRange
0.00%	3.0600us	6	510ns	325ns	737ns	cudaGetLastError
0.00%	3.0510us	3	1.0170us	931ns	1.1630us	cuInit
0.00%	2.0790us	3	693ns	621ns	779ns	cuDriverGetVersion
0.00%	1.9120us	1	1.9120us	1.9120us	1.9120us	cudaUnbindTexture
0.00%	1.1120us	1	1.1120us	1.1120us	1.1120us	cudaGetDeviceCount

The profile displays two parts of time. First is the time consumed by each kernel. For m1.2 most time is consumed by *cuda::detail::implicit\_convolve\_sgemm* this kernel(37.01% 50.022ms). The second is the time consumed by each API calls like *cudaFree* *cudaMemcpy*. For m1.2, *cudaStreamCreateWithFlags* API call costed 46.76%(1.85341s, called 18 times) of total time of all API calls.

## Milestone 2: A New CPU Layer in MXNet

### M2.1 Add CPU forward implementation

#### M2.1.1 Description of implementation

Batch size: B

Input features/channels: C inputs ( $H \times W$ ).

Convolution Layer: M filters ( $K \times K$ ).

Output Features/channels: M outputs  $(H - K + 1) \times (W - K + 1)$ .

We use 'C' to represent the number of input feature maps and use 'H' to represent the height of each input map image, and the width of each is 'W'. Assume that the input feature maps are stored in a 3D array X [B,C, H, W]. We have M outputs feature map, and each size is  $(H - K + 1) \times (W - K + 1)$ . The following shows a sequential code of CNN for forward propagation path.

Our Code here:

```
const int B = x.shape_[0];
const int M = y.shape_[1];
const int C = x.shape_[1];
const int H = x.shape_[2];
const int W = x.shape_[3];
const int K = k.shape_[3];
int H_out = H - K + 1;
int W_out = W - K + 1;
for (int b = 0; b < B; ++b) {
    //CHECK_EQ(0, 1) << "Missing an ECE408 CPU implementation!";
    /* ... a bunch of nested loops later...
        y[b][m][h][w] += x[b][c][h + p][w + q] * k[m][c][p][q];
    */
    for (int m = 0; m < M; m++) // for each output feature map
        for (int h = 0; h < H_out; h++) // for each output element
            for (int w = 0; w < W_out; w++) {
```

```

y[b][m][h][w] = 0;
for(int c = 0; c < C; c++)    // sum over all input feature maps
    for(int p = 0; p < K; p++)    // KxK filter
        for(int q = 0; q < K; q++)
            y[b][m][h][w] += x[b][c][h + p][w + q] * k[m][c][p][q];

```

## M2.1.2 Result and performance

Our baseline cpu implementation correctness and performance results is:

\* Running python m2.1.py

New Inference

Loading fashion-mnist data... done

Loading model... done

Op Time: 9.045332

Correctness: 0.8562 Model: ece408-high

## Milestone 3

---

### 3.1 Add a simple GPU forward implementation

In this milestone, we implemented the basic GPU forward convolution and then used shared memory to load filter kernel and tiled input as an optimization.

Shared memory size :

filter kernel  $C * K * K * \text{sizeof}(\text{float})$

input:  $(\text{TILE\_WIDTH} + K - 1)^2$

### 3.2 Create a GPU profile with nvprof.

\* Running nvprof python m3.1.py

New Inference

Loading fashion-mnist data... done

==311== NVPROF is profiling process 311, command: python m3.1.py

Loading model... done

Op Time: 0.103691

Correctness: 0.8562 Model: ece408-high

==311== Profiling application: python m3.1.py

==311==

Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
34.28%	65.149ms	1	65.149ms	65.149ms	65.149ms	mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)
20.65%	39.252ms	1	39.252ms	39.252ms	39.252ms	sgemm_sm35_ldg_tn_128x8x256x16x32
10.21%	19.399ms	1	19.399ms	19.399ms	19.399ms	void mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>, mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul, mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)
10.20%	19.394ms	2	9.6968ms	461.21us	18.932ms	void cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>(cudnnTensorStruct, float const *, cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>, cudnnTensorStruct*, float, cudnnTensorStruct*, int, cudnnTensorStruct*)
9.49%	18.029ms	3	6.0095ms	3.1360us	17.535ms	[CUDA memcpy DtoD]
7.63%	14.495ms	1	14.495ms	14.495ms	14.495ms	void cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)
4.19%	7.9578ms	13	612.13us	1.8560us	5.5119ms	[CUDA memcpy HtoD]
1.92%	3.6513ms	1	3.6513ms	3.6513ms	3.6513ms	sgemm_sm35_ldg_tn_64x16x128x8x32
0.59%	1.1199ms	1	1.1199ms	1.1199ms	1.1199ms	void mshadow::cuda::SoftmaxKernel<int=8, float, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>>(mshadow::gpu, int=2, unsigned int)
0.40%	754.10us	12	62.841us	2.0800us	380.99us	void mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.23%	437.72us	2	218.86us	17.408us	420.31us	void mshadow::cuda::MapPlanKernel<mshadow::sv::plusto, int=8, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::Broadcast1DExp<mshadow::Tensor<mshadow::gpu, int=1, float>, float, int=2, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.21%	394.01us	1	394.01us	394.01us	394.01us	sgemm_sm35_ldg_tn_32x16x64x8x16
0.01%	23.264us	1	23.264us	23.264us	23.264us	void mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::ReduceWithAxisExp<mshadow::red::maximum, mshadow::Tensor<mshadow::gpu, int=3, float>, float, int=3, bool=1, int=2>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.01%	9.9520us	1	9.9520us	9.9520us	9.9520us	[CUDA memcpy DtoH]

==311== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
46.50%	2.82762s	18	157.09ms	19.678us	1.42871s	cudaStreamCreateWithFlags
29.30%	1.78217s	13	137.09ms	1.4970us	514.67ms	cudaFree
20.66%	1.25635s	23	54.624ms	280.94us	1.24862s	cudaMemGetInfo
1.36%	82.565ms	1	82.565ms	82.565ms	82.565ms	cudaDeviceSynchronize
1.29%	78.531ms	25	3.1412ms	6.6570us	42.555ms	cudaStreamSynchronize
0.56%	33.779ms	44	767.70us	11.330us	20.586ms	cudaMalloc
0.23%	13.735ms	8	1.7169ms	14.264us	5.7998ms	cudaMemcpy2DAsync
0.03%	1.7250ms	4	431.25us	37.675us	1.5485ms	cudaStreamCreate
0.03%	1.5263ms	4	381.57us	378.31us	385.95us	cuDeviceTotalMem
0.02%	1.2081ms	24	50.338us	16.392us	357.39us	cudaLaunch
0.02%	1.1688ms	352	3.3200us	307ns	115.66us	cuDeviceGetAttribute
0.01%	330.49us	9	36.720us	21.596us	85.209us	cudaMemcpy
0.00%	172.89us	114	1.5160us	925ns	4.3730us	cudaEventCreateWithFlags
0.00%	116.52us	4	29.130us	21.645us	33.854us	cuDeviceGetName

0.00%	113.49us	30	3.7820us	854ns	11.621us	cudaSetDevice
0.00%	87.898us	104	845ns	592ns	2.0750us	cudaDeviceGetAttribute
0.00%	82.803us	146	567ns	346ns	1.5660us	cudaSetupArgument
0.00%	52.197us	2	26.098us	21.625us	30.572us	cudaStreamCreateWithPriority
0.00%	42.172us	24	1.7570us	682ns	5.1630us	cudaConfigureCall
0.00%	26.656us	10	2.6650us	1.3070us	8.4980us	cudaGetDevice
0.00%	13.254us	17	779ns	613ns	1.0220us	cudaPeekAtLastError
0.00%	9.2620us	2	4.6310us	2.1300us	7.1320us	cudaEventRecord
0.00%	7.1010us	1	7.1010us	7.1010us	7.1010us	cudaStreamGetPriority
0.00%	5.8340us	6	972ns	485ns	2.1170us	cuDeviceGetCount
0.00%	5.7600us	2	2.8800us	2.2420us	3.5180us	cudaStreamWaitEvent
0.00%	4.6220us	6	770ns	471ns	1.2630us	cuDeviceGet
0.00%	4.3260us	5	865ns	667ns	1.0680us	cudaGetLastError
0.00%	4.2820us	2	2.1410us	2.0100us	2.2720us	cudaDeviceGetStreamPriorityRange
0.00%	3.7080us	3	1.2360us	1.0900us	1.4200us	cuDriverGetVersion
0.00%	3.4300us	3	1.1430us	1.0180us	1.3220us	cuInit
0.00%	1.5140us	1	1.5140us	1.5140us	1.5140us	cudaGetDeviceCount

## Contribution

Chaohua Shang: Write and run the code, write most of the milestone 1 and milestone 3 part in report

Haojia: Write and run the code, assist Chaohua and Jiayue with the report by providing information and writing part of it.

Jiayue Wang: Write and run the code, write most of the milestone 2 part in report