

ECE408 Report

APLUSPLUS

1. Baseline Results

Milestone 1 : Get Started

M1.1: Run the Baseline Forward Pass

Run CPU code in rai

Our Results as below(add elapsed time for whole m1.1.py program):

* Running time python m1.1.py

New Inference

Loading fashion-mnist data... done

Loading model... done

EvalMetric: {'accuracy': 0.8673}

11.40user 11.48system

0:21.23elapsed 107%CPU (0avgtext+0a

vgdata 1628360maxresident)k

0inputs+2624outp

uts (0major+27159minor)pagefau

lts 0swaps

M1.2/1.3: Run the Baseline GPU implementation and generate a NVPROF profile

Modified rai_build.yml according to introduction.

Mxnet GPU layer performance results is as below:

* Running nvprof python m1.2.py

New Inference

Loading fashion-mnist data... done

==310== NVPROF is profiling process 310, command: python m1.2.py

Loading model...[02:36:12] src/operator/././cudnn_algoreg-inl.h:112: Running performance tests to find the best convolution algorithm, this can take a while... (setting env variable MXNET_CUDNN_AUTOTUNE_DEFAULT to 0 to disable)
done

EvalMetric: {'accuracy': 0.8673}

==310== Profiling application: python m1.2.py

==310== Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
37.01%	50.022ms	1	50.022ms	50.022ms	50.022ms	void cudnn::detail::implicit_convolve_sgemm<float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>(int, int, int, float const *, int, cudnn::detail::implicit_convolve_sgemm<float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>*, float const *, kernel_conv_params, int, float, float, int, float const *, float const *, int, int)
28.67%	38.751ms	1	38.751ms	38.751ms	38.751ms	sgemm_sm35_ldg_tn_128x8x256x16x32
14.34%	19.385ms	2	9.6923ms	459.06us	18.926ms	void cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>(cudnnTensorStruct, float const *, cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>, cudnnTensorStruct*, float, cudnnTensorStruct*, int, cudnnTensorStruct*)
10.69%	14.445ms	1	14.445ms	14.445ms	14.445ms	void cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)
4.53%	6.1185ms	13	470.66us	1.5360us	4.1914ms	[CUDA memcpy HtoD]
2.75%	3.7160ms	1	3.7160ms	3.7160ms	3.7160ms	sgemm_sm35_ldg_tn_64x16x128x8x32
0.82%	1.1115ms	1	1.1115ms	1.1115ms	1.1115ms	void mshadow::cuda::SoftmaxKernel<int=8, float, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>>(mshadow::gpu, int=2, unsigned int)
0.55%	748.43us	12	62.369us	2.0800us	378.04us	void mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.32%	433.43us	2	216.72us	16.639us	416.79us	void mshadow::cuda::MapPlanKernel<mshadow::sv::plusto, int=8, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::Broadcast1DExp<mshadow::Tensor<mshadow::gpu, int=1, float>, float, int=2, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.29%	390.68us	1	390.68us	390.68us	390.68us	sgemm_sm35_ldg_tn_32x16x64x8x16
0.02%	22.399us	1	22.399us	22.399us	22.399us	void mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::ReduceWithAxisExp<mshadow::red::maximum, mshadow::Tensor<mshadow::gpu, int=3, float>, float, int=3, bool=1, int=2>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.01%	10.016us	1	10.016us	10.016us	10.016us	[CUDA memcpy DtoH]

==310== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
46.76%	1.85341s	18	102.97ms	17.717us	926.36ms	cudaStreamCreateWithFlags
28.75%	1.13936s	10	113.94ms	696ns	322.52ms	cudaFree
20.66%	818.70ms	24	34.112ms	230.81us	811.60ms	cudaMemGetInfo
3.24%	128.32ms	25	5.1329ms	5.5140us	83.313ms	cudaStreamSynchronize
0.31%	12.278ms	8	1.5348ms	8.4420us	6.1684ms	cudaMemcpy2DAsync
0.17%	6.5406ms	42	155.73us	10.600us	1.1805ms	cudaMalloc
0.03%	1.3602ms	4	340.06us	338.87us	342.68us	cuDeviceTotalMem
0.02%	862.41us	114	7.5640us	625ns	307.12us	cudaEventCreateWithFlags
0.02%	848.24us	352	2.4090us	248ns	63.433us	cuDeviceGetAttribute
0.01%	527.45us	23	22.932us	11.237us	89.276us	cudaLaunch
0.01%	419.09us	6	69.847us	59.623us	81.884us	cudaMemcpy
0.01%	233.20us	4	58.298us	35.204us	83.353us	cudaStreamCreate
0.00%	102.25us	4	25.561us	17.241us	32.322us	cuDeviceGetName
0.00%	78.556us	32	2.4540us	695ns	8.6440us	cudaSetDevice
0.00%	70.258us	110	638ns	418ns	2.3920us	cudaDeviceGetAttribute
0.00%	57.402us	147	390ns	253ns	1.1690us	cudaSetupArgument
0.00%	40.865us	2	20.432us	18.308us	22.557us	cudaStreamCreateWithPriority
0.00%	26.085us	23	1.1340us	492ns	3.3300us	cudaConfigureCall
0.00%	24.680us	10	2.4680us	1.3040us	5.9290us	cudaGetDevice

0.00%	9.0930us	1	9.0930us	9.0930us	9.0930us	cudaBindTexture
0.00%	8.7240us	16	545ns	407ns	891ns	cudaPeekAtLastError
0.00%	4.3310us	6	721ns	240ns	1.4230us	cuDeviceGetCount
0.00%	4.1760us	2	2.0880us	1.5150us	2.6610us	cudaStreamWaitEvent
0.00%	3.7180us	1	3.7180us	3.7180us	3.7180us	cudaStreamGetPriority
0.00%	3.4330us	6	572ns	415ns	874ns	cuDeviceGet
0.00%	3.4270us	2	1.7130us	1.3330us	2.0940us	cudaEventRecord
0.00%	3.2360us	2	1.6180us	1.4250us	1.8110us	cudaDeviceGetStreamPriorityRange
0.00%	3.0600us	6	510ns	325ns	737ns	cudaGetLastError
0.00%	3.0510us	3	1.0170us	931ns	1.1630us	cuInit
0.00%	2.0790us	3	693ns	621ns	779ns	cuDriverGetVersion
0.00%	1.9120us	1	1.9120us	1.9120us	1.9120us	cudaUnbindTexture
0.00%	1.1120us	1	1.1120us	1.1120us	1.1120us	cudaGetDeviceCount

The profile displays two parts of time. First is the time consumed by each kernel. For m1.2 most time is consumed by *cudnn::detail::implicit_convolve_sgemm* this kernel(37.01% 50.022ms). The second is the time consumed by each API calls like *cudaFree* *cudaMemcpy*. For m1.2, *cudaStreamCreateWithFlags* API call costed 46.76%(1.85341s, called 18 times) of total time of all API calls.

Milestone 2: A New CPU Layer in MXNet

M2.1 Add CPU forward implementation

M2.1.1 Description of implementation

Batch size: B

Input features/channels: C inputs (H × W).

Convolution Layer: M filters (K × K).

Output Features/channels: M outputs (H − K + 1) × (W − K + 1).

We use 'C' to represent the number of input feature maps and use 'H' to represent the height of each input map image, and the width of each is 'W'. Assume that the input feature maps are stored in a 3D array X [B, C, H, W]. We have M outputs feature map, and each size is (H − K + 1) × (W − K + 1). The following shows a sequential code of CNN for forward propagation path.

Our Code here:

```
const int B = x.shape_[0];
const int M = y.shape_[1];
const int C = x.shape_[1];
const int H = x.shape_[2];
const int W = x.shape_[3];
const int K = k.shape_[3];
int H_out = H - K + 1;
int W_out = W - K + 1;
for (int b = 0; b < B; ++b) {
    //CHECK_EQ(0, 1) << "Missing an ECE408 CPU implementation!";
```

```

/* ... a bunch of nested loops later...
  y[b][m][h][w] += x[b][c][h + p][w + q] * k[m][c][p][q];
*/
for(int m = 0; m < M; m++)          // for each output feature map
  for(int h = 0; h < H_out; h++)      // for each output element
    for(int w = 0; w < W_out; w++) {
      y[b][m][h][w] = 0;
      for(int c = 0; c < C; c++)      // sum over all input feature maps
        for(int p = 0; p < K; p++)    // KxK filter
          for(int q = 0; q < K; q++)
            y[b][m][h][w] += x[b][c][h + p][w + q] * k[m][c][p][q];
    }

```

M2.1.2 Result and performance

Our baseline cpu implementation correctness and performance results is:

* Running python m2.1.py

New Inference

Loading fashion-mnist data... done

Loading model... done

Op Time: 9.045332

Correctness: 0.8562 Model: ece408-high

Milestone 3

3.1 Add a simple GPU forward implementation

In this milestone, we implemented the basic GPU forward convolution and then used shared memory to load filter kernel and tiled input as an optimization.

Shared memory size :

filter kernel $C * K * K * \text{sizeof(float)}$

input: $(\text{TILE_WIDTH} + K - 1)^2$

3.2 Create a GPU profile with nvprof.

* Running nvprof python m3.1.py

New Inference

Loading fashion-mnist data... done

==311== NVPROF is profiling process 311, command: python m3.1.py

Loading model... done

Op Time: 0.103691

Correctness: 0.8562 Model: ece408-high

==311== Profiling application: python m3.1.py

==311==

Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
34.28%	65.149ms	1	65.149ms	65.149ms	65.149ms	mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)
20.65%	39.252ms	1	39.252ms	39.252ms	39.252ms	sgemm_sm35_ldg_tn_128x8x256x16x32
10.21%	19.399ms	1	19.399ms	19.399ms	19.399ms	void mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>, mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul, mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)
10.20%	19.394ms	2	9.6968ms	461.21us	18.932ms	void cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>(cudnnTensorStruct, float const *, cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>, cudnnTensorStruct*, float, cudnnTensorStruct*, int, cudnnTensorStruct*)
9.49%	18.029ms	3	6.0095ms	3.1360us	17.535ms	[CUDA memcpy DtoD]
7.63%	14.495ms	1	14.495ms	14.495ms	14.495ms	void cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)
4.19%	7.9578ms	13	612.13us	1.8560us	5.5119ms	[CUDA memcpy HtoD]
1.92%	3.6513ms	1	3.6513ms	3.6513ms	3.6513ms	sgemm_sm35_ldg_tn_64x16x128x8x32
0.59%	1.1199ms	1	1.1199ms	1.1199ms	1.1199ms	void mshadow::cuda::SoftmaxKernel<int=8, float, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>>(mshadow::gpu, int=2, unsigned int)
0.40%	754.10us	12	62.841us	2.0800us	380.99us	void mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.23%	437.72us	2	218.86us	17.408us	420.31us	void mshadow::cuda::MapPlanKernel<mshadow::sv::plusto, int=8, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::Broadcast1DExp<mshadow::Tensor<mshadow::gpu, int=1, float>, float, int=2, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.21%	394.01us	1	394.01us	394.01us	394.01us	sgemm_sm35_ldg_tn_32x16x64x8x16
0.01%	23.264us	1	23.264us	23.264us	23.264us	void mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan<mshadow::expr::ReduceWithAxisExp<mshadow::red::maximum, mshadow::Tensor<mshadow::gpu, int=3, float>, float, int=3, bool=1, int=2>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.01%	9.9520us	1	9.9520us	9.9520us	9.9520us	[CUDA memcpy DtoH]

==311== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
46.50%	2.82762s	18	157.09ms	19.678us	1.42871s	cudaStreamCreateWithFlags
29.30%	1.78217s	13	137.09ms	1.4970us	514.67ms	cudaFree
20.66%	1.25635s	23	54.624ms	280.94us	1.24862s	cudaMemGetInfo
1.36%	82.565ms	1	82.565ms	82.565ms	82.565ms	cudaDeviceSynchronize
1.29%	78.531ms	25	3.1412ms	6.6570us	42.555ms	cudaStreamSynchronize
0.56%	33.779ms	44	767.70us	11.330us	20.586ms	cudaMalloc
0.23%	13.735ms	8	1.7169ms	14.264us	5.7998ms	cudaMemcpy2DAsync
0.03%	1.7250ms	4	431.25us	37.675us	1.5485ms	cudaStreamCreate

0.03%	1.5263ms	4	381.57us	378.31us	385.95us	cuDeviceTotalMem
0.02%	1.2081ms	24	50.338us	16.392us	357.39us	cudaLaunch
0.02%	1.1688ms	352	3.3200us	307ns	115.66us	cuDeviceGetAttribute
0.01%	330.49us	9	36.720us	21.596us	85.209us	cudaMemcpy
0.00%	172.89us	114	1.5160us	925ns	4.3730us	cudaEventCreateWithFlags
0.00%	116.52us	4	29.130us	21.645us	33.854us	cuDeviceGetName
0.00%	113.49us	30	3.7820us	854ns	11.621us	cudaSetDevice
0.00%	87.898us	104	845ns	592ns	2.0750us	cudaDeviceGetAttribute
0.00%	82.803us	146	567ns	346ns	1.5660us	cudaSetupArgument
0.00%	52.197us	2	26.098us	21.625us	30.572us	cudaStreamCreateWithPriority
0.00%	42.172us	24	1.7570us	682ns	5.1630us	cudaConfigureCall
0.00%	26.656us	10	2.6650us	1.3070us	8.4980us	cudaGetDevice
0.00%	13.254us	17	779ns	613ns	1.0220us	cudaPeekAtLastError
0.00%	9.2620us	2	4.6310us	2.1300us	7.1320us	cudaEventRecord
0.00%	7.1010us	1	7.1010us	7.1010us	7.1010us	cudaStreamGetPriority
0.00%	5.8340us	6	972ns	485ns	2.1170us	cuDeviceGetCount
0.00%	5.7600us	2	2.8800us	2.2420us	3.5180us	cudaStreamWaitEvent
0.00%	4.6220us	6	770ns	471ns	1.2630us	cuDeviceGet
0.00%	4.3260us	5	865ns	667ns	1.0680us	cudaGetLastError
0.00%	4.2820us	2	2.1410us	2.0100us	2.2720us	cudaDeviceGetStreamPriorityRange
0.00%	3.7080us	3	1.2360us	1.0900us	1.4200us	cuDriverGetVersion
0.00%	3.4300us	3	1.1430us	1.0180us	1.3220us	culnit
0.00%	1.5140us	1	1.5140us	1.5140us	1.5140us	cudaGetDeviceCount

2.Optimization Approach and Results

1,2,3.how you identified the optimization opportunity

We use following optimised functions:

Unrolled the input features and convert the convolution to **matrix multiplication**.

While optimizing the matrix multiplication we considered the memory access coalesce, control divergence and global memory access. Used tiled matrix multiplication and applied some modification according to our input data size.

First of all, we implemented the MP3. if we use 32 as TILE_WIDTH, every warp will have control divergence while loading W and X into shared memory then, and warp near the boundary will have control divergency writing the product into Y according as the input size of w is 50*25, and X size is 25*576. So, we chose the 25 as the TILE_WIDTH. In this case, threads have no control divergence while loading input W into shared memory, but it still have control divergency while loading X and writing Y. After changed the size of TILE_WIDTH, the time decreased from 0.150975s to 0.140642s.

We also tried constant memory. We thought it could be effective by using constant memory to improve the performance a little bit because filter W is used by every batch

and w. However, the results turned out to be opposite. The time it spent was 0.19625s, which was longer than the time using shared memory 0.140642s.

We then analyzed the tiled matrix multiplication. If we using 25 as Tile width, the kernel will launch 2×24 thread blocks for each input batch . During the first iteration, block(0,0), block(0,1)... block(0,23) loaded $W_{0,0} \dots W_{24,24}$ (25X25) into shared memory respectively. So, the kernel access the same W elements multiple times in order to calculate Y in one batch. We figured that we could use one thread to calculate two output by using one block to load two block size of X and calculate two block size of Y. In this case, we only launched 2×12 thread blocks for one batch. And block(0,0) loaded $W_{0,0} \dots W_{24,24}$ (25X25) and $X_{0,0} \dots X_{24,24}$ (25X25) $X_{25,0} \dots X_{49,49}$ (25X25), and calculated $X_{0,0} \dots X_{24,49}$ (25X50). This reduced 1X global memory access. After implemented above, the time decreased to 0.099665s.

Similarly, we can use one thread to calculate four output. Block(0,0) will load $W_{0,0} \dots W_{49,24}$ (50X25) and $X_{0,0} \dots X_{49,49}$ (25X50) and calculate $X_{0,0} \dots X_{49,49}$ (50X50). This will reduce the number of parallel computation but will reduce the time in the global memory access. The time performance of this method is 0.071211s. The time of Forward_kernel which is matrix multiplication kernel is 43.79ms

New Inference

Loading fashion-mnist data... done

==315== NVPROF is profiling process 315, command: python final.py

Loading model... done

Op Time: 0.071211

Correctness: 0.8562 Model: ece408-high

==315== Profiling application: python final.py

==315== Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
25.21%	43.790ms	1	43.790ms	43.790ms	43.790ms	mxnet::op::forward_kernel(float*, float*, float*, int, int, int, int, int, int)

- 4, Any external references used during identification or development of the optimization
 - [1]Guangming Tan, Linchuan Li, Sean Treichler, Everett Phillips, Yungang Bao, and Ninghui Sun. Fast implementation of DGEMM on Fermi GPU. In Supercomputing 2011, SC '11, pages 35:1–35:11, New York, NY, USA, 2011. ACM.
 - [2] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In NIPS, pages 1097–1105, 2012.

5, How your team organized and divided up this work.

In Final stage:

We worked together to come up with the optimized function in order to accelerate the model.

6, References (as needed)

[1]NVIDIA cuDNN - GPU accelerated deep learning.

[2]Chapter 16 - 3rd-Edition-Chapter16-case-study-DNN-FINAL-corrected.pdf

7,(Optional) Suggestions for Improving Next Year

We hope to make our project in flexible topics. For example,Game Tree Search, CoMD GPU implementation. And use criteria to grade so that student could Play creativity.

Contribution

Chaohua Shang: Write and run the code, wrote most of the milestone 1, milestone 3 and final part in report

Haojia: Write and run the code, wrote the final part and assist Chaohua and Jiayue with the report by providing information and writing part of it.

Jiayue Wang: Write and run the code, write most of the milestone 2 and final part in report